

AI Safety Framework Coding Codebook

Revision of 12 August 2026 includes a 'Quick test' for Users in the very first table (section 1. Authoritative code list) and removes the mutual exclusivity for C09-C16. Some quick tests are included in the Codes themselves.

Canonical version 8 — structural revisions, 11 August 2026

Supersedes Canonical v7 (7 August 2026). This revision is structural: it fixes identifiers, cross-references, field placement and output conventions, and adds one inclusion criterion to C14. It does not change what any other code covers.

1. Authoritative code list

This table is authoritative. Where anything elsewhere in this document — including the revision notes — implies a different code set or a different total, this table governs.

Fourteen active codes: six change-family (C01, C03–C07) and eight content-family (C08, C09, C11–C16).

Each code has a **permanent ID** that never changes and a **display number** that reflects its position in this document. Cite codes by ID. Display numbers have already moved twice (7 August 2026) and may move again if codes are retired; IDs will not. Retired IDs are tombstoned and are never reissued, so a dataset coded under any version maps forward one-to-one.

ID	Canonical name	Family	Display no.	Direction	Quick test
C01	Competitor/Conditionality Clause Introduction	Change	Code 1	Yes	Rival's actions ease our own commitment?
C02	Threshold Tightening/Loosening — RETIRED 7 Aug 2026, merged into C04	—	—	—	
C03	Threshold Existence Change	Change	Code 2	Yes	Does the gate exist at all here?
C04	Obligation & Threshold Strength Shift	Change	Code 3	Yes + facet	Gate persists — bar or firmness moved?
C05	Adding or Dropping a Risk Theme/Domain	Change	Code 4	Yes	Category's presence changed between versions?
C06	Definitional Drift	Change	Code 5	Yes	Same category — what counts narrowed/broadened?
C07	Governance Tightening/Loosening	Change	Code 6	Yes	Governance body gained/lost authority?

ID	Canonical name	Family	Display no.	Direction	Quick test
C08	Frontier AI Risk Definition & Framing	Content	Code 7	Yes (multi)	Defines, divides, or motivates frontier risk?
C09	Risk Severity Levels & Thresholds	Content	Code 8	No	States what the threshold is?
C10	Scope of Risks — RETIRED 7 Aug 2026	—	—	—	
C11	Autonomy/Loss of Control Risk Theme	Content	Code 9	No	Independent action or evading control?
C12	Cyber-related Risk Theme	Content	Code 10	No	About cyberattack or cyber defense?
C13	CBRN Risk Theme	Content	Code 11	No	About chemical/bio/radiological/nuclear threats?
C14	Governance, Transparency and Disclosure	Content	Code 12	No	About oversight, disclosure, or delegation?
C15	Persuasion and Influence Risk Theme	Content	Code 13	No	About manipulating belief or behavior?
C16	Evaluation and Risk Management Practices	Content	Code 14	No	Describes HOW risk is assessed/managed?

Canonical names. Use the name exactly as written above. Earlier versions varied ("Adding/Dropping" for C05, "&" or "static Governance" for C14, "Persuasion/Manipulation" for C15); those variants are not canonical and should not appear in output, column headers, or cross-references.

2. How to tag a passage

Assignment is binary and unordered. For each passage, assign every code whose criteria the passage satisfies. There is no limit and no ranking — no Tag 1, no Tag 2, no judgment about which code a passage is "most centrally" about.

Most passages will satisfy one or two codes. Some will satisfy three or more. Some will satisfy none, and an untagged passage is a valid result — not every sentence needs a code.

Where a code carries a direction, record it alongside the assignment (see §3.2). A code assigned without its direction is incomplete.

Worked examples

"CBRN remains a tracked category in both versions, but the definition of what counts as CBRN uplift has narrowed."

- C13 CBRN Risk Theme
- C06 Definitional Drift [narrowed]

OpenAI's 2025 line: "Persuasion category risks do not fit the criteria for inclusion."

- C05 Adding or Dropping a Risk Theme/Domain [dropped] — Persuasion was tracked in 2023 and is not in 2025
- C15 Persuasion and Influence Risk Theme — the passage is still substantively about persuasion risk

C10 Scope of Risks was retired on 7 August 2026; see the tombstone before C11 for where in/out-of-scope language now goes.

A governance passage that both tightens oversight and explains why:

- C07 Governance Tightening/Loosening [tightened]
- C14 Governance, Transparency and Disclosure

A passage can need three or more codes when three or more things are genuinely true of it — see the DeepMind example under C14. That is uncommon but correct when it happens.

3. Coding conventions

3.1 "Skip" means assign 0

Several exclusion criteria below use the word **skip** (e.g. "→ skip; note separately"). Throughout this codebook, *skip* means **do not assign that code to this passage**. It never means excluding the passage from the dataset, dropping the row, or leaving a field empty. Where the study records assignments as a binary matrix, *skip* means enter 0 in that code's column.

3.2 Direction vocabularies

Where a code carries a direction, the value must come from that code's closed list below. Do not invent values and do not write free text. Codes not listed here carry no direction.

ID	Code	Permitted values
C01	Competitor/Conditionality Clause Introduction	introduced · expanded
C03	Threshold Existence Change	introduced · removed · reintroduced · architecture_replaced
C04	Obligation & Threshold Strength Shift	Direction: tightened · loosened. Facet (required): modality · bar · both. See §3.3.
C05	Adding or Dropping a Risk Theme/Domain	added · dropped · reintroduced · split · merged
C06	Definitional Drift	narrowed · broadened
C07	Governance Tightening/Loosening	tightened · loosened

ID	Code	Permitted values
C08	Frontier AI Risk Definition & Framing	Multi-select, one or more of: A-umbrella · A-framing · A-motivation
C09, C11–C16	(content codes)	No direction. Record NA.

architecture_replaced (C03) covers a whole measurement or gating architecture being retired and replaced by a structurally different one — the case the code describes in its inclusion criteria. split and merged (C05) cover a category dividing into several or several consolidating into one, as when 2023's single Model Autonomy category became AI Self-improvement, Long-range Autonomy and Autonomous Replication and Adaptation in 2025.

3.3 C04 facet rule

C04 covers two facets — a shift in **verb modality** and a shift in a **numeric or stated bar**. Record which facet moved:

- **modality** — the verb register moved and the stated bar did not.
- **bar** — the stated trigger value moved and the verb register did not.
- **both** — both moved, in the same direction.

Divergent case. If the two move in opposite directions (the language firms up while the bar loosens, or the reverse), record the direction of the **bar**, set facet to both, and flag the passage as ambiguous with the reason "modality and bar diverge." Do not split the passage into two assignments.

3.4 Consistency rules

These follow from the exclusion criteria in the entries below. They are stated here so they can be checked directly.

Rule	Basis
C03 and C04 are mutually exclusive	Reciprocal exclusions, reinforced by the two Quick tests: a gate either exists or does not (C03), or persists while its bar or firmness moves (C04).
C06 is mutually exclusive with C04	C06 excludes threshold-value change; C04 excludes substantive change in what counts.
C06 is mutually exclusive with C05	C06 excludes a category's presence changing; C05 excludes a definitional shift within a persisting category.
C07 requires C14	C07 is a change-family supplement to governance content; a governance-robustness shift presupposes a governance passage.
C03 and C05 may co-occur	Stated in C03's example (the Persuasion drop touches both).
C08 and C16 may co-occur	Explicit dual-tag instruction in both entries, for passages that both frame why ongoing risk-scanning matters and describe the scanning practice.
C08 and C09 — check	Keep both only where the passage both fixes or divides the risk concept and states a graded tier. If only one is true, keep only that code.

Rule	Basis
C09 and C16 may co-occur	No longer mutually exclusive (struck 11 August 2026). Assign both where a passage both states what a threshold is AND describes how it is measured, evaluated, or revised. Assign C09 only where the passage states the threshold with no accompanying practice/methodology content; assign C16 only where the passage describes practice with no threshold content.
C11, C12, C13, C15 co-occur freely	No exclusions operate between risk themes, or between a risk theme and any other code.

3.5 Change-family codes require two versions

C01 and C03–C07 are change-family: each requires a prior version of the same framework to compare against. Where no prior counterpart exists for a passage, these six codes are all 0. A commitment appearing for the first time is not, by itself, a change to anything.

4. The codes

C01 (Code 1). Competitor/Conditionality Clause Introduction

Code name	Competitor/Conditionality Clause Introduction
ID	C01
Family	Change
Theme	Conditional Weakening
Definition	Passages that introduce or expand language conditioning a lab's safety commitment on what peers or competitors do — i.e. a clause allowing the lab's own commitments (safeguard stringency, deployment timing, or release schedule) to ease if a rival ships a comparable capability without equivalent precautions.
Inclusion criteria	Explicit reference to competitor or rival behaviour as a condition on the lab's own commitment; phrases like "if a comparable model is released without similar safeguards," "in light of the competitive landscape." Also covers deployment-timing or release-schedule adjustments conditioned on competitor behaviour (e.g. "we may accelerate deployment timelines if a comparable model has already been released") — the same underlying move (a safety-relevant commitment becoming conditional on a rival's actions) applied to a different lever than safeguard stringency.
Exclusion criteria	Excludes general aspirational softening that does not specifically reference competitor behaviour.
Direction	introduced · expanded
Example	OpenAI's 2025 v2 Preparedness Framework adding a clause about easing safeguards if a rival ships something comparable without them. (For OpenAI, these are found in 2025.) Note: this instance is the safeguard-stringency flavour of the code. The passage does not address deployment timing or release schedules, but these may appear in other labs' frameworks in relation to competitor developments. A clause conditioning timing or release schedule on competitor behaviour or market developments would equally qualify.
Counter-example	"We aim to apply these safeguards where practicable."

Retired: C02 — Threshold Tightening/Loosening

Retired 7 August 2026. Merged into C04 Obligation & Threshold Strength Shift. The two were never cleanly separable: the codebook's own worked examples for both were the same scenario (a pause commitment softening), and every clean pilot instance was either a verb-modality shift (C04's territory) or a substantive scope change already covered by C06 Definitional Drift. Passages about a threshold's bar value changing between versions now go to C04.

The ID C02 is permanently retired and is never reissued.

C03 (Code 2). Threshold Existence Change

Code name	Threshold Existence Change
ID	C03

Family	Change
Theme	Threshold Existence
Definition	Quick test: does the threshold or gate exist at all in this version? (appeared, vanished, or was replaced by a structurally different mechanism.) Passages where a threshold or trigger point appears for the first time with no prior equivalent, disappears entirely with no replacement, or reappears in a later version after being absent. MERGED CODE: combines the former "Threshold Existence Change (newly introduced/appear/disappear)" and "Threshold Introduction/Removal" entries, which the pilot review (20 July 2026) identified as duplicates.
Inclusion criteria	A threshold or category present in one version and simply absent in another, with no reworded replacement standing in its place; a new introduction; a reintroduction after absence; a removal. Also covers an entire measurement or gating architecture being retired and replaced by a structurally different one (e.g. an overall post-mitigation Scorecard system replaced by per-category capability thresholds) — the prior system's threshold disappears even though a new gating mechanism appears in its place; record this as direction architecture_replaced. Also covers a tracked level threshold being downgraded so that it no longer functions as a Tracked-category gate (e.g. folded into a Research Category) — assign C05 Adding or Dropping a Risk Theme/Domain alongside. Language such as "introduces," "newly tracked," "for the first time," "reintroduced," "restored," or explicit notice of a prior removal.
Exclusion criteria	Excludes cases where a threshold persists under the same or a directly successor name but only its stringency, numeric value, or wording changes — that is C04 Obligation & Threshold Strength Shift.
Direction	introduced · removed · reintroduced · architecture_replaced
Example	OpenAI's Preparedness Framework dropping Persuasion in its 2025 version touches both this code and C05 Adding or Dropping a Risk Theme/Domain — assign both. Pilot example: OpenAI 2023's overall post-mitigation Scorecard gate ("medium or below to deploy") being replaced in 2025 by per-category High/Critical capability thresholds, coded C03 [architecture_replaced].
Counter-example	"CBRN remains a tracked category in both versions, but the definition of what counts as CBRN uplift has narrowed." (This is C06 Definitional Drift, not existence change — the category never left, and it is the substantive meaning of what counts that shifted, not a bar value or verb modality.)

C04 (Code 3). Obligation & Threshold Strength Shift

Code name	Obligation & Threshold Strength Shift
ID	C04
Family	Change
Theme	Commitment Firmness & Threshold Strength
Definition	Quick test: the threshold still exists in both versions — did its bar or its firmness move? Quick test for facet: did the verb change (modality), the number change (bar), or both? MERGED CODE (merged with the former C02 Threshold Tightening/Loosening on 7 August 2026). Passages tracking whether a commitment's verb modality shifts between mandatory

	("will," "shall," "must"), conditional ("will, unless X"), and aspirational ("aim to," "intend to") register across versions, becoming firmer or more vague — OR where a specific numeric bar or trigger value for the same named threshold moves between versions (e.g. ">2x time saved" becoming ">5x time saved"), tightening or loosening the same metric. These two facets often arrive bundled in the same passage. Assign the code once and record which facet moved (modality, bar, or both) per §3.3; do not split the passage into two assignments.
Inclusion criteria	Modality facet. Any version-to-version comparison where the same commitment's verb strength changes. Signals: "will" becoming "will aim to" or vice versa; "will trigger" becoming "may trigger"; "at our discretion," "provided that," "unless," "except where," "as appropriate given the circumstances." Bar facet. The same named threshold's stated trigger value changing between versions while the threshold itself persists. Signals: a quantity, percentage, multiplier, time window, count, or named benchmark score that appears in both versions against the same trigger and differs between them (e.g. ">2x time saved" becoming ">5x time saved"; "within 24 hours" becoming "within 7 days"). The trigger must persist under the same or a directly successor name — a new threshold appearing, or an architecture being restructured, is C03.
Exclusion criteria	Excludes the competitor-specific case already captured under C01 Competitor/Conditionality Clause Introduction. Excludes a threshold or category stated for the first time with no prior version to compare, or an entire measurement architecture being restructured — that is C03 Threshold Existence Change. Excludes a substantive change in WHAT counts toward a persisting threshold's definition (e.g. a capability that used to qualify no longer does, or vice versa) where the trigger's own wording or value does not change — that is C06 Definitional Drift.
Direction	tightened · loosened. Facet (required): modality · bar · both. For the divergent case, see §3.3.
Example	DeepMind's v1.0 stating it "will pause deployment" upon a Critical Capability Level being reached, versus v3.1 stating it "will evaluate appropriate mitigations" upon the same trigger. [loosened, facet: modality] — Anthropic's v1.0 pause commitment reading as a firm stop at a set capability level, versus v3.3 leaving more room for judgment about timing and competitors. [loosened, facet: modality]
Coder note	On the DeepMind example: a first reading may suggest [tightened] because the later wording is more elaborate. It is a loosening — the commitment moved from a definite action ("will pause") to a discretionary assessment ("will evaluate appropriate mitigations"). Judge the direction by the firmness of what the lab is bound to do, not by the specificity of the language describing it.
Counter-example	"We will publish our framework annually," appearing identically worded across versions — no shift in either modality or bar value.

READ ME — telling three easily-confused codes apart

C08 Frontier AI Risk Definition & Framing, C05 Adding or Dropping a Risk Theme/Domain, and C06 Definitional Drift are frequently confused, since all three touch how a risk concept or category is described. Ask these questions, in order:

- **C08 Frontier AI Risk Definition & Framing** — what does the lab say frontier, existential or catastrophic risk MEANS, or how does it carve risk into buckets, within a single version?
- **C05 Adding or Dropping a Risk Theme/Domain** — did a whole category's PRESENCE change between versions: tracked in one, gone (or reintroduced) in another?

- **C06 Definitional Drift** — does the substantive definition of an already-present risk concept narrow or broaden between versions, even though the category itself persists in both?

Worked example: OpenAI's 2025 line "Persuasion category risks do not fit the criteria for inclusion" is assigned C05 [dropped] (Persuasion was tracked in 2023 and is not in 2025) plus C15 Persuasion and Influence Risk Theme (the passage is still substantively about persuasion). If instead CBRN stayed tracked in both versions but its definition of qualifying uplift narrowed, that would be C06, not C05 — the category never left.

C05 (Code 4). Adding or Dropping a Risk Theme/Domain

Code name	Adding or Dropping a Risk Theme/Domain
ID	C05
Family	Change
Theme	Scope Shift — risk themes being added, dropped, or reintroduced
Definition	Passages where an entire risk category, domain, or tracked capability area is added, removed, or reintroduced between versions, as opposed to a threshold changing within a category that persists.
Inclusion criteria	A named category present in one version and absent in another, or a category that disappears in one version and reappears in a later one; language like "no longer tracked," "new category," "reintroduced." Also covers a single category dividing into several (split) or several consolidating into one (merged).
Exclusion criteria	Excludes cases where a category's definition shifts but the category itself remains present across all versions — that is C06 Definitional Drift.
Direction	added · dropped · reintroduced · split · merged
Example	OpenAI's Preparedness Framework dropping the Persuasion category between the 2023 Beta and v2 [dropped]; DeepMind's Frontier Safety Framework cutting the Autonomy domain and later folding misalignment work into a single domain in v3.1 [merged]. Pilot example: 2023's single "Model Autonomy" category becoming AI Self-improvement (Tracked), Long-range Autonomy (Research) and Autonomous Replication and Adaptation (Research) in 2025 [split].
Counter-example	"CBRN remains a tracked category in both versions, but the definition of what counts as CBRN uplift has narrowed."

C06 (Code 5). Definitional Drift

Code name	Definitional Drift
ID	C06
Family	Change
Theme	Substantive Narrowing or Broadening of an Already-Present Risk Concept
Definition	Added following pilot review, 20 July 2026, in response to a question about whether a risk concept defined differently across two versions was already covered elsewhere — it was not. Passages where the substantive definition of a risk concept or category narrows or broadens between two versions of the same framework, even though the concept or

	category itself is present in both versions and no threshold or category presence changed. Change-family only: requires a prior version to compare against.
Inclusion criteria	The umbrella definition of frontier, catastrophic or existential risk changing in substance between versions (e.g. what counts as "severe harm" expanding or contracting). A named risk theme (e.g. CBRN) remaining tracked in both versions, but what qualifies as an instance of that risk narrowing or broadening (e.g. what counts as "CBRN uplift"). Record a direction: [narrowed] = the definition covers less than before; [broadened] = the definition covers more than before. The narrowing/broadening test: can you name one concrete real-world example that would have qualified under the old wording but would not qualify under the new wording, or vice versa? If yes, that is Definitional Drift. If you can only point to different words describing the same real-world cases either way, it is stylistic — see Exclusion criteria.
Exclusion criteria	Excludes a category's presence changing (tracked → gone, or vice versa) — that is C05 Adding or Dropping a Risk Theme/Domain. Excludes a numeric or tiered threshold changing while the underlying definition of the risk stays the same — that is C04 Obligation & Threshold Strength Shift. Excludes a risk being moved inside or outside document scope, or handed to another instrument: route that per the tombstone note preceding C11 — to the relevant risk-theme code, to C14 Governance, Transparency and Disclosure, or to C16 Evaluation and Risk Management Practices as appropriate. Excludes cases where the definition is worded identically or with only stylistic changes across versions; use the narrowing/broadening test to check.
Direction	narrowed · broadened
Example	A. "CBRN remains a tracked category in both versions, but the definition of what counts as CBRN uplift has narrowed." [narrowed] — assign C13 CBRN Risk Theme alongside. B. A closer call, showing the test in action: a Cybersecurity Critical threshold reading "exploits... across all software projects" in one version and "exploits... in many hardened real-world critical systems" in the next. At first glance this could be read as rewording. Applying the test: can you name something that would have counted before but not now? Yes — a small hobbyist app or non-critical internal tool would have been "software" under the old wording but would not obviously be a "hardened real-world critical system" under the new one. That is a real narrowing, not style. [narrowed]
Counter-example	OpenAI's Preparedness Framework dropping the Persuasion category outright between 2023 and 2025 — the category's presence changed, not its definition. That is C05. Also a counter-example for the test itself: two versions that swap "will identify security vulnerabilities" for "will discover security flaws" — no real-world case would qualify under one wording and not the other, so the test says skip it as stylistic (assign 0; see §3.1).

C07 (Code 6). Governance Tightening/Loosening

Code name	Governance Tightening/Loosening
ID	C07
Family	Change
Theme	Governance Robustness Change
Definition	Added following pilot review, 20 July 2026. Passages where the internal governance apparatus overseeing safety decisions — committees, reporting lines, approval and veto authority — is strengthened or weakened between versions of the same framework. Change-family only: requires a prior version to compare against. C07 is rare — if

	assigning it more than a handful of times, re-verify each instance names a specific body, reporting line, or authority; general safety-culture language is not C07.
Inclusion criteria	A body gaining or losing oversight or veto authority (e.g. a Board committee newly given power to "reverse a decision"); a reporting line added or cut (e.g. less direct reporting between Audit and the Board); a committee created, dissolved, or merged; escalation or emergency fast-track procedures added or removed. Record a direction: [tightened] = more oversight, more layers, more veto power, more direct reporting; [loosened] = less oversight, fewer layers, less veto power, less direct reporting.
Exclusion criteria	Excludes a governance body or process described identically across both versions — assign C14 Governance, Transparency and Disclosure only, with no direction. Excludes the stated reason for a change, which stays under C14's explanatory-language criterion. Excludes cases with no prior version to compare — that is C14 only, not a shift.
Direction	tightened · loosened
Dependency	C07 requires C14 on the same passage (see §3.4).
Example	2025 OpenAI SAG passage adding "OpenAI Leadership can approve or reject these recommendations, and our Board's Safety and Security Committee provides oversight of these decisions" — an approval and oversight layer not present in the 2023 SAG description. [tightened]
Counter-example	SAG described in near-identical terms in both versions, with no change to who reports to whom or who holds final say — no direction applies; C14 only.

READ ME — telling three easily-confused codes apart (repeated for convenience)

C08 Frontier AI Risk Definition & Framing, C05 Adding or Dropping a Risk Theme/Domain, and C06 Definitional Drift are frequently confused, since all three touch how a risk concept or category is described. Ask these questions, in order:

- **C08 Frontier AI Risk Definition & Framing** — what does the lab say frontier, existential or catastrophic risk MEANS, or how does it carve risk into buckets, within a single version?
- **C05 Adding or Dropping a Risk Theme/Domain** — did a whole category's PRESENCE change between versions: tracked in one, gone (or reintroduced) in another?
- **C06 Definitional Drift** — does the substantive definition of an already-present risk concept narrow or broaden between versions, even though the category itself persists in both?

Worked example: OpenAI's 2025 line "Persuasion category risks do not fit the criteria for inclusion" is assigned C05 [dropped] (Persuasion was tracked in 2023 and is not in 2025) plus C15 Persuasion and Influence Risk Theme (the passage is still substantively about persuasion). If instead CBRN stayed tracked in both versions but its definition of qualifying uplift narrowed, that would be C06, not C05 — the category never left.

C08 (Code 7). Frontier AI Risk Definition & Framing

Code name	Frontier AI Risk Definition & Framing
ID	C08
Family	Content

Theme	Definition of concepts such as frontier risk
Definition	Passages that tell you what the lab takes AI frontier risk to BE, or why it matters — either by defining the big concept, dividing catastrophic or existential risk into major kinds, or motivating why frontier or catastrophic risk requires attention (most typically, a misuse-type versus misalignment-type split, in each lab's own vocabulary).
Inclusion criteria	[A-umbrella] Defining the big concept. Sentences fixing what catastrophic, frontier or existential risk means, or "unknown unknowns" that may form part of frontier risk. Signals: "we mean," "refers to," "focus on," glossary entries. A sentence without these verbs still counts if it pins down the meaning. [A-framing] Dividing frontier or catastrophic risk into major kinds (e.g. CBRN, misuse, loss of autonomy). If a lab defines frontier risk and also names other risk categories or themes (e.g. CBRN), assign the relevant risk-theme code as well. [A-motivation] Motivating why frontier or catastrophic risk matters, or why the lab is investing in studying or tracking it — opening statements about the state of the field, the stakes involved, or the case for taking the risk seriously, including framing sentences that precede or introduce a fuller definition. Signals: "we believe," "it is critical that," "fallen short of where we need to be." A passage may carry more than one of the three; record all that apply.
Exclusion criteria	Dictionary entry for ONE category ("CBRN: risks of models assisting in... a CBRN threat") → assign the relevant risk-category code only; this code applies only when the passage structures frontier risk as a whole. Definition of a tool, not a risk ("A Capability Threshold is a prespecified level of AI capability...") → skip (assign 0; see §3.1); note separately. Mission or value talk with no connection to frontier or catastrophic risk at all ("we aim to build beneficial AI for all of humanity") → skip. Mission or value talk that motivates why frontier or catastrophic risk specifically matters, or why the lab studies or tracks it, is covered under [A-motivation] rather than skipped. Note: lab-specific proxies or benchmarks used to render a shared risk are not coded — handled later in analytic memos. Clarifying note (added 7 August 2026). A bare label naming the "Unknown Unknowns" bucket stays under this code, being definitional and taxonomic; the ongoing process of monitoring for and identifying new risk categories — the scanning itself — is assigned to C16 Evaluation and Risk Management Practices, because it is a practice or methodology question ("how you do it"). Where a passage does both, framing why ongoing scanning matters AND describing the scanning practice itself, assign both this code and C16.
Direction	Multi-select: A-umbrella · A-framing · A-motivation
Example	"'Catastrophic risk' as used in our RSP refers generally to risks of the most severe potential harms from advanced AI, such as existential threats or fundamental destabilization of global systems." (Anthropic) [A-umbrella]. "We identify CCLs for two kinds of risks: misuse risk and risks related to machine learning R&D and misalignment." (Google) [A-framing]. "We believe the scientific study of catastrophic risks from AI has fallen far short of where we need to be." (OpenAI 2023) [A-motivation].
Counter-example	"Exact scenarios of loss of control risks are speculative and difficult to precisely specify." (xAI) — talks about defining without defining. Field note only.

C09 (Code 8). Risk Severity Levels & Thresholds

Code name	Risk Severity Levels & Thresholds
------------------	-----------------------------------

ID	C09
Family	Content
Theme	Category and level regulation — what capability levels trigger High or Critical, and why
Definition	Passages that GRADE a risk: how severe, at what tier, past what line. Rule of thumb — this code captures what the threshold IS, not how it gets measured, revised, or governed. For evaluation methodology, benchmark results, mitigation practice and emerging-risk monitoring, see C16 Evaluation and Risk Management Practices.
Inclusion criteria	Tiers and their content: CBRN-3, CCL and TCL levels, Low–Medium–High–Critical tables, xAI's three behaviour buckets. Numeric cutoffs: "above X% on benchmark Y." No-go, halt or deployment-blocking triggers: "capability X would trigger a halt on deployment Y" → assign the relevant risk-theme code alongside this one.
Exclusion criteria	Dividing risk into kinds without grading it ("misuse risk and... misalignment risk") → C08 Frontier AI Risk Definition & Framing. Governance structures built on the tiers (approval boards, reporting lines) → skip; consider C14. Passages that describe ONLY the procedure for measuring or updating thresholds, with no statement of what the threshold itself is (e.g. "we may adjust alert thresholds...", "we plan to add additional thresholds...") → skip this code (assign 0; see §3.1) and assign C16 instead. Where a passage both states a threshold's content and describes how it is measured, evaluated, or revised, assign both this code and C16 — they are not mutually exclusive (revised 11 August 2026; see §3.4).
Direction	None. Record NA.
Example	"CBRN uplift level 1: Provides low to medium resourced actors uplift in reference scenarios resulting in additional expected harm at severe scale." (Google) — a tier and its content.
Counter-example	"We may run early warning evaluations more frequently or adjust alert thresholds if the rate of progress suggests our safety buffer is no longer adequate." (Google) — a procedure for revising a threshold, not what the threshold is.

Retired: C10 — Scope of Risks

Retired 7 August 2026. Scope of Risks (INSIDE / OUTSIDE / DEFERRED / DELEGATED) was dropped as its own code: in practice its [excluded] and [deferred] passages always duplicated C05 Adding or Dropping a Risk Theme/Domain or C03 Threshold Existence Change, and the code was a frequent source of confusion.

Passages stating that a named risk is in scope, out of scope, deferred, or delegated elsewhere now route as follows. If the passage is substantively about a specific risk theme (C11 Autonomy/Loss of Control, C12 Cyber-related, C13 CBRN, or C15 Persuasion and Influence), assign that risk-theme code. If it hands responsibility for a risk to another instrument, policy or body — the delegation case — assign C14 Governance, Transparency and Disclosure, which carries an explicit delegation criterion for this purpose. If it is about the practice of monitoring, evaluating or managing that risk, assign C16 Evaluation and Risk Management Practices. If none genuinely fits, it is correct to leave the passage untagged.

The ID C10 is permanently retired and is never reissued.

C11 (Code 9). Autonomy/Loss of Control Risk Theme

Code name	Autonomy/Loss of Control Risk Theme
ID	C11
Family	Content
Theme	Ability to act independently, improve, and/or evade developer control
Definition	Model autonomy relates to the level of independent decision-making and action a model can take without human intervention. Risks stem from a model's ability to autonomously accelerate its own capability development, potentially compromising or escaping human control. Passages where the document discusses: an AI system's capacity to execute complex, long-horizon, uninstructed actions that are fully independent; operating without human intervention, enabling scaled misuse; a model developing autonomy and evading developer control; or a model independently acquiring adaptive resources (renting its own cloud servers, etc.), self-exfiltration capabilities, or autonomous replication and adaptation (surviving, making copies of itself).
Inclusion criteria	Language like "progression of autonomy," "loss of control," enabling "actors to run scaled misuse," "self-initiated actions," "fully autonomous agents" capable of "long-term / long-range / long time-horizon planning," model "adaptation" to complex scenarios; intentional model underperformance during checks ("sandbagging"), "autonomous replication," "completes chained actions," initiating high-risk tasks without human instruction, an "autonomous model creating uplift."
Exclusion criteria	Excludes simple triggers that are baseline-level rule-based automation. Note that "autonomy" is adjacent to but separate from "agency" (e.g. defined by AWS as the "scope of actions that an AI model is enabled/permitted to take within an environment, and how much a human bounds an agent's capabilities/permissions").
Direction	None. Record NA.
Example	"In conjunction with a long range autonomy capability, models that could bypass OpenAI's technical safeguards constraining model activity, such as sandboxing or monitoring services, could compromise OpenAI's ability to..."
Counter-example	"The model solves well-specified tasks demonstrating an understanding of programming." (Excluded — medium risk; no demonstration of adaptation or long-term contextual understanding; no ability to replicate or self-exfiltrate; no lack of human controllability; model not adapting to a human attempting to shut it down.)

C12 (Code 10). Cyber-related Risk Theme

Code name	Cyber-related Risk Theme
ID	C12
Family	Content
Theme	Cyber operations, cyber threats and cybersecurity
Definition	The ability of an AI model to assist in the development of tools and executing operations for cyber defence and cyber offence; enterprise security posture and operations related to cyber risk identification, mitigation, and novel threats. Passages where the document discusses: risks and criticality pertaining to the use of a model by malicious actors to compromise digital systems with intent to cause harm; the potential that a frontier model may enable a novice or moderately skilled actor to discover and exploit high-value security

	vulnerabilities (resulting in loss or casualties) or offensive cybersecurity methods; or to independently instruct or automatically deploy cyber operations against a hardened target (see also C11 Autonomy/Loss of Control Risk Theme).
Inclusion criteria	Language like "cyber threat and automated threat intelligence," "AI-orchestrated cyber espionage campaign," "kill chain," "cyber tactics, techniques, procedures," "offensive cyber operations," "ransomware," "IP theft," "scaled spear-phishing," "adaptive malware creation able to reliably evade commonly used detection solutions," thresholds for "API deployments," "cyberdefense" in relation to model weight security, "cyberoffense," "cyber red teaming," "scaled cyber attack," "threat modeling exercises and analysis of the current fraud and scams landscape."
Exclusion criteria	Excludes simple or generic statements about resilience, which can pertain to cyber as much as to other risk categories or themes.
Direction	None. Record NA.
Example	"The Offensive Cyber Operations Threshold focuses on the potential that a frontier model may provide material uplift in excess of other publicly available research or existing tools..."
Counter-example	A generic statement that includes an allusion to "typical scams," but in fact relates to another risk theme: "This makes models a significant aid to biased journalism, get-out-the-vote campaigns, and typical scams..."

C13 (Code 11). CBRN Risk Theme

Code name	CBRN Risk Theme
ID	C13
Family	Content
Theme	Ability to accelerate access to biological, chemical, radiological and/or nuclear research knowledge; discovery and R&D related to CBRN and dual-use hazards; lowering the barrier to entry to design or deploy dangerous agents
Definition	Risks of models assisting in the development, preparation, and/or execution of a chemical, biological, radiological or nuclear ("CBRN") threat. This includes a model's ability to assist with lab work or synthesis and production of certain agents, to genetically modify or cultivate agents, to design or deploy devices, and to disseminate this knowledge, or to access fissile materials or ballistics.
Inclusion criteria	Passages where the document discusses the ideation, acquisition, magnification, formulation and release of chemical or biological threats; potential for biological, chemical or nuclear terror events; pathogens, radiological devices, agents; references to high mortality or ease of transmission, health risk, deaths and disruption; or considerations on CBRN severity and criticality. Language like "biological weapon," "biosecurity," "biological threat vectors," "CDC," "lab automation," radioactive material acquisition, models providing "expert-level, interactive instruction that provides material uplift" related to lab processes, drug discovery or weapons design, and red-teaming in relation to chemical, biological, radiological or nuclear threats. Note: a CBRN passage can genuinely need more than one code — as a risk theme, and also as C05 Adding or Dropping a Risk Theme/Domain, and also as a narrowing or broadening under C06. Assign all that genuinely apply.
Exclusion criteria	Excludes references not scoped to a Tracked or Research Category by virtue of not requiring security controls or safeguards against misuse. For example: model output to human prompts (say, on human anatomical processes) providing a user increasing levels of

	detail, with no High or Critical capability thresholds, no associated risk of severe harm, and no increased likelihood of biological or chemical terror.
Direction	None. Record NA.
Example	"The model can enable an expert to develop a highly dangerous novel threat vector (e.g. comparable to a novel CDC Class A biological agent...)"
Counter-example	"CBRN risk threshold rationale presenting minimal change from status quo where other resources (e.g. textbooks, search engines) already exist."

C14 (Code 12). Governance, Transparency and Disclosure

Code name	Governance, Transparency and Disclosure
ID	C14
Family	Content
Theme	Senior executive and management committee and board governance; transparency and disclosure to the public; enterprise risk governance as a governance practice; corporate governance and disclosure, including independent oversight
Definition	Passages where the document discusses internal governance practices, a governance function, centralised review processes and decision-making (senior advisory groups and committees), accountability and reporting protocols, Board committees, and third-party audit.
Inclusion criteria	Language like "a Board's Safety and Security Committee providing oversight of these decisions," "internal transparency" documented through reporting to a Senior Advisory Group, the Board having the option to "reverse a decision and/or mandate a revised course of action," transparency in material findings being shared with internal stakeholders. Explanatory language. Any stated reason offered for a change flagged under any change-family code — "we revised this because..." "in response to..." — including disclosing, or no longer disclosing, a changed threshold to the Board or a benchmark to the public. Delegation (added 11 August 2026, on retirement of C10 Scope of Risks). Passages assigning responsibility for a named risk, or for its oversight, to an instrument, policy, body or process outside the framework — e.g. "our Usage Policy sets forth our standards for...", "this is addressed under our Model Spec," "handled through our trust and safety processes," "governed by our AI Principles." Delegation is an allocation-of-responsibility decision and belongs here. The target instrument need not be named specifically; a passage stating that a risk is handled elsewhere qualifies even where "elsewhere" is unspecified. Where the passage is also substantively about a named risk theme, assign that theme code alongside. Note. If the passage also describes the governance apparatus itself becoming more or less robust relative to a prior version, assign C07 Governance Tightening/Loosening as well, with a direction.
Exclusion criteria	Excludes simple compliance requirements linked to an enterprise process (e.g. access management practices), such as "ensure all access is regularly audited and reviewed." This pertains to infosecurity best practice rather than high-level governance; consider C16 Evaluation and Risk Management Practices. Excludes a sentence on expectations linked to capability thresholds, which could be assigned C09 Risk Severity Levels & Thresholds rather than this code, and C04 Obligation & Threshold Strength Shift depending on context. For example: "As we conduct further

	research into these and other risk domains, we expect these CCLs to evolve and for several CCLs at higher levels or in other risk domains to be added." Excludes a sentence describing how mitigation plans might be developed or controls might change in relation to critical capability levels, e.g.: "As we better understand the risks posed by models at different CCLs, and the contexts in which our models will be deployed, we will develop mitigation plans that map the CCLs to the security and deployment levels described." This could be assigned C09 (or C04 depending on version changes), and not this code. A mitigation plan is, at its most abstract, an element of good governance practice; nonetheless this code does not include mitigation plans. Where a broader phrasing or paragraph describes mitigation plans or revisions linked to controls, thresholds, capabilities or the rank-ordering of risk themes, code it through that lens — C09, or C04, possibly with a second code.
Direction	None. Record NA. (For version-to-version governance change, see C07.)
Example	A. "The Safety and Security Committee (SSC) of the OpenAI Board of Directors (Board) will be given visibility into processes, and can review decisions and otherwise require reports and information from OpenAI Leadership as necessary to fulfill the Board's oversight role. Where necessary, the Board may reverse a decision and/or mandate a revised course of action." B. DeepMind explaining that Harmful Manipulation was reintroduced "in response to emerging evidence of manipulative model behavior." This passage carries three codes: C05 Adding or Dropping a Risk Theme/Domain [reintroduced] (the category came back), C15 Persuasion and Influence Risk Theme (the topic), and C14 (the stated reason for reintroducing it). All three are true of the passage.
Counter-example	"Ensure all access is regularly audited and reviewed." This is about internal access parameters — infosecurity best practice for users — not corporate governance as such, even though it refers to auditing and to review, reporting and disclosure. Consider C16.

C15 (Code 13). Persuasion and Influence Risk Theme

Code name	Persuasion and Influence Risk Theme
ID	C15
Family	Content
Theme	Strategic distortion of human belief or behaviour, including deception and manipulation
Definition	Passages where the document discusses influencing or convincing people to alter core beliefs or to act on model-generated content, including by exploiting emotional or cognitive vulnerabilities; or AI systems deceiving humans, concealing their true capabilities or intentions, or providing false information. "Persuasive capability" is defined via the level of human persuasive content that can be commoditised by a model (OpenAI): risks stemming from the strategic distortion of human behaviour. For example, a model that can "create (potentially interactive) content with persuasive effectiveness strong enough to convince almost anyone to take action on a belief that goes against their natural interest" (Critical risk level). Also covers passages where persuasion, harmful or advanced persuasion, or belief manipulation relates to complex topics (e.g. an election) or conduct (e.g. self-harm); to safety and bias-prevention challenges, especially in the context of influence operations; or where a human's susceptibility or vulnerability may be tactically targeted (e.g. urging a depressed person to self-harm). A passage under this code may also carry C11 Autonomy/Loss of Control Risk Theme, but need not. Note: influence operations and public sentiment manipulation may also be tracked

	outside frontier risk policies, through product-abuse or trust and safety frameworks, or within cyber-related risk assessments (footnote as applicable).
Inclusion criteria	Language like "persuasive threat," "coordinated influence campaigns," a weapon "for extracting secrets and interfering with democracy," persuasive "sycophancy" (in some cases), "perception hacking," "generative disinformation," "generate deceptive content at scale." Language like "deceive," "mislead," or "conceal capabilities" in some cases.
Exclusion criteria	Excludes simple accuracy issues or hallucination, since those are not intentional deception.
Direction	None. Record NA.
Example	A model hits a dangerous tipping point or capability level if it demonstrates the ability to execute "scalable, high-stakes interactive manipulation."
Counter-example	"The model sometimes generates factually incorrect outputs." (Excluded — this is an error, not intentional deception.)

C16 (Code 14). Evaluation and Risk Management Practices

Code name	Evaluation and Risk Management Practices
ID	C16
Family	Content
Theme	Evaluation methodology and risk management practice
Definition	Passages describing HOW risk is assessed, tested, monitored or mitigated in practice — evaluation methodology, capability elicitation technique, benchmark and performance results (e.g. model versus human baseline, X% versus Y%), safety-baseline procedures, mitigation and safeguard implementation practice, generic security-control best practices, and the ongoing process of scanning for and capturing emerging or "unknown unknown" risk categories — as distinct from passages that define WHAT threshold or severity level triggers a designation.
Inclusion criteria	Evaluation approach descriptions (e.g. scalable evaluations, deep dives, red-teaming cadence, capability elicitation methodology). Illustrative or reported benchmark and performance results, including comparisons of model versus human baseline performance. Mitigation or safeguard implementation practices not themselves tied to naming a specific severity trigger. Generic security-control best practices (access management, patching, audits, defence in depth) where not already captured by another code — assess C16 after all other codes have been decided. The emerging or unknown-risk capture process: the ongoing work of monitoring for, identifying and assessing whether a new or nascent risk category warrants tracking, as distinct from a named "Unknown Unknowns" category label itself, which is C08.
Exclusion criteria	Statements that ONLY name or define a specific High, Critical, Low or Medium threshold or numeric cutoff, with no accompanying description of measurement, evaluation, or revision practice → C09 Risk Severity Levels & Thresholds only. Where a passage does both — states the threshold's content AND describes how it is measured, evaluated, or revised — assign both C09 and this code; they are not mutually exclusive (revised 11 August 2026; see §3.4).

	Governance body, committee, reporting-line or approval-authority content → C14 Governance, Transparency and Disclosure instead. A bare label naming a risk-classification bucket (e.g. “Unknown Unknowns,” “Tracked Categories,” “Research Categories”) without describing the ongoing monitoring process → C08 Frontier AI Risk Definition & Framing instead. Clarifying note (added 7 August 2026). A bare label naming the “Unknown Unknowns” bucket stays with C08, being definitional and taxonomic; the ongoing process of monitoring for and identifying new risk categories — the scanning itself — is assigned here, because it is a practice or methodology question (“how you do it”). Where a passage does both, framing why ongoing scanning matters AND describing the scanning practice itself, assign both C08 and this code.
Direction	None. Record NA.
Example	"GPT vs search eval: Post-PhD professionals trained in biology (specifically genetics, virology, or microbiology) experience no improvement in success rates for writing CDC Class B biological agent development instructions using a search engine and GPT versus just a search engine (X% vs X%)..." (OpenAI 2023) — a reported benchmark result, not a threshold definition. "We invest deeply in developing or adopting new science-backed evaluations that provide high precision and high recall indications of whether a covered system has reached a capability threshold..." (OpenAI 2025) — evaluation methodology. "The Preparedness team conducts research, evaluations, monitoring, and forecasting of risks... monitoring for unknown unknowns and making the case for inclusion... of any new risk categories as they emerge." (OpenAI 2023) — emerging-risk capture process.
Counter-example	"[High] The model's impact is equivalent to giving every OpenAI researcher a highly performant mid-career research engineer assistant, relative to those researchers' 2024 baseline." (OpenAI 2025) — this states what the threshold IS → C09, not this code, even though it also involves a comparative benchmark framing.

